

demo-mnist

April 22, 2025

```
[12]: import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import RMSprop
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt
from sklearn import metrics

# Load the OCR dataset

# The MNIST dataset is a built-in dataset provided by Keras.
# It consists of 70,000 28x28 grayscale images, each of which displays a single
    ↳ handwritten digit from 0 to 9.
# The training set consists of 60,000 images, while the test set has 10,000
    ↳ images.

(x_train, y_train), (x_test, y_test) = mnist.load_data()

# X_train and X_test are our array of images while y_train and y_test are our
    ↳ array of labels for each image.
# The first tuple contains the training set features (X_train) and the training
    ↳ set labels (y_train).
# The second tuple contains the testing set features (X_test) and the testing
    ↳ set labels (y_test).
# For example, if the image shows a handwritten 7, then the label will be the
    ↳ integer 7.

plt.imshow(x_train[0], cmap='gray') # imshow() function which simply displays
    ↳ an image.
plt.show() # cmap is responsible for mapping a specific colormap to the values
    ↳ found in the array that you passed as the first argument.


# image appears black and white and that each axis of the plot ranges from 0 to
    ↳ 28.

# This is because of the format that all the images in the dataset have:
```

```
# 1. All the images are grayscale, meaning they only contain black, white and
    ↳grey.
# 2. The images are 28 pixels by 25 pixels in size (28x28).

print(x_train[0])

# image data is just an array of digits. You can almost make out a 5 from the
    ↳pattern of the digits in the array.
# Array of 28 values
# a grayscale pixel is stored as a digit between 0 and 255 where 0 is black,
    ↳255 is white and values in between are different shades of gray.
# Therefore, each value in the [28][28] array tells the computer which color to
    ↳put in that position when we display the actual image.
```



```
[ [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0]
  [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0]
  [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0]
  [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0]
```

```

[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 3 18 18 18 126 136
175 26 166 255 247 127 0 0 0 0]
[ 0 0 0 0 0 0 0 0 30 36 94 154 170 253 253 253 253 253
225 172 253 242 195 64 0 0 0 0]
[ 0 0 0 0 0 0 0 49 238 253 253 253 253 253 253 253 253 251
93 82 82 56 39 0 0 0 0 0]
[ 0 0 0 0 0 0 0 18 219 253 253 253 253 253 198 182 247 241
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 80 156 107 253 253 205 11 0 43 154
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 14 1 154 253 90 0 0 0 0
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 139 253 190 2 0 0 0
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 11 190 253 70 0 0 0
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 35 241 225 160 108 1
0 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 81 240 253 253 119
25 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 45 186 253 253
150 27 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 16 93 252
253 187 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 249
253 249 64 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 46 130 183 253
253 207 2 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 39 148 229 253 253 253
250 182 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 24 114 221 253 253 253 253 201
78 0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 23 66 213 253 253 253 253 198 81 2
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 18 171 219 253 253 253 253 195 80 9 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 55 172 226 253 253 253 253 244 133 11 0 0 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 136 253 253 253 212 135 132 16 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0]
[ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0]

```

```
[13]: # reformat our X_train array and our X_test array because they do not have the
      ↪ correct shape.
      # Reshape the data to fit the model
      print("X_train shape", x_train.shape)
      print("y_train shape", y_train.shape)
      print("X_test shape", x_test.shape)
      print("y_test shape", y_test.shape)

      # Here you can see that for the training sets we have 60,000 elements and the
      ↪ testing sets have 10,000 elements.

      # y_train and y_test only have 1 dimensional shapes because they are just the
      ↪ labels of each element.

      # x_train and x_test have 3 dimensional shapes because they have a width and
      ↪ height (28x28 pixels) for each element.

      # (60000, 28, 28) 1st parameter in the tuple shows us how much image we have
      ↪ 2nd and 3rd parameters are the pixel values from x to y (28x28)
      # The pixel value varies between 0 to 255.
      # (60000,) Training labels with integers from 0-9 with dtype of uint8. It has
      ↪ the shape (60000,).
      # (10000, 28, 28) Testing data that consists of grayscale images. It has the
      ↪ shape (10000, 28, 28) and the dtype of uint8. The pixel value varies between
      ↪ 0 to 255.
      # (10000,) Testing labels that consist of integers from 0-9 with dtype uint8.
      ↪ It has the shape (10000,).
```

```
X_train shape (60000, 28, 28)
y_train shape (60000,)
X_test shape (10000, 28, 28)
y_test shape (10000,)
```

```
[14]: # X: Training data of shape (n_samples, n_features)
      # y: Training label values of shape (n_samples, n_labels)
      # 2D array of height and width, 28 pixels by 28 pixels will just become 784
      ↪ pixels (28 squared).
      # Remember that X_train has 60,000 elements, each with 784 total pixels so
      ↪ will become shape (60000, 784).
      # Whereas X_test has 10,000 elements, each with each with 784 total pixels so
      ↪ will become shape (10000, 784).

      x_train = x_train.reshape(60000, 784)
      x_test = x_test.reshape(10000, 784)
```

```

x_train = x_train.astype('float32') # use 32-bit precision when training a
↳neural network, so at one point the training data will have to be converted
↳to 32 bit floats. Since the dataset fits easily in RAM, we might as well
↳convert to float immediately.
x_test = x_test.astype('float32')
x_train /= 255 # Each image has Intensity from 0 to 255
x_test /= 255

# Regarding the division by 255, this is the maximum value of a byte (the input
↳feature's type before the conversion to float32),
# so this will ensure that the input features are scaled between 0.0 and 1.0.
# USING svm-https://mgta.gmu.edu/courses/ml-with-python/
↳handwrittenDigitRecognition.php#:~:
↳text=Remember%20that%20X_train%20has%2060%2C000,keras.

```

```

[15]: # Convert class vectors to binary class matrices
num_classes = 10
y_train = np.eye(num_classes)[y_train] # Return a 2-D array with ones on the
↳diagonal and zeros elsewhere.
y_test = np.eye(num_classes)[y_test] # f your particular categories is present
↳then it mark as 1 else 0 in remain row

```

```

[16]: # Define the model architecture
model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,))) # The input_shape
↳argument is passed to the foremost layer. It comprises of a tuple shape,
model.add(Dropout(0.2)) # DROP OUT RATIO 20%
model.add(Dense(512, activation='relu')) #returns a sequence of vectors of
↳dimension 512
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

```

```

[17]: # Compile the model
model.compile(loss='categorical_crossentropy', # for a multi-class
↳classification problem
              optimizer=RMSprop(),
              metrics=['accuracy'])

```

```

[18]: # Train the model
batch_size = 128 # batch_size argument is passed to the layer to define a batch
↳size for the inputs.
epochs = 20
history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,

```

```

        verbose=1, # verbose=1 will show you an animated progress
    bar eg. [=====]
        validation_data=(x_test, y_test)) # Using validation_data
    means you are providing the training set and validation set yourself,
        # validation_split means
    you only provide a training set and keras splits it into a training set and
    a validation set

```

Epoch 1/20

469/469 [=====] - 12s 24ms/step - loss: 0.2539 - accuracy: 0.9216 - val_loss: 0.1443 - val_accuracy: 0.9547

Epoch 2/20

469/469 [=====] - 9s 20ms/step - loss: 0.1039 - accuracy: 0.9665 - val_loss: 0.0730 - val_accuracy: 0.9789

Epoch 3/20

469/469 [=====] - 10s 22ms/step - loss: 0.0742 - accuracy: 0.9767 - val_loss: 0.0693 - val_accuracy: 0.9795

Epoch 4/20

469/469 [=====] - 11s 22ms/step - loss: 0.0565 - accuracy: 0.9827 - val_loss: 0.0669 - val_accuracy: 0.9811

Epoch 5/20

469/469 [=====] - 10s 22ms/step - loss: 0.0480 - accuracy: 0.9850 - val_loss: 0.0747 - val_accuracy: 0.9790

Epoch 6/20

469/469 [=====] - 10s 22ms/step - loss: 0.0396 - accuracy: 0.9872 - val_loss: 0.0699 - val_accuracy: 0.9817

Epoch 7/20

469/469 [=====] - 10s 20ms/step - loss: 0.0342 - accuracy: 0.9889 - val_loss: 0.0687 - val_accuracy: 0.9833

Epoch 8/20

469/469 [=====] - 11s 22ms/step - loss: 0.0294 - accuracy: 0.9903 - val_loss: 0.0661 - val_accuracy: 0.9836

Epoch 9/20

469/469 [=====] - 10s 22ms/step - loss: 0.0255 - accuracy: 0.9923 - val_loss: 0.0745 - val_accuracy: 0.9822

Epoch 10/20

469/469 [=====] - 10s 22ms/step - loss: 0.0220 - accuracy: 0.9930 - val_loss: 0.0734 - val_accuracy: 0.9836

Epoch 11/20

469/469 [=====] - 9s 19ms/step - loss: 0.0202 - accuracy: 0.9934 - val_loss: 0.0679 - val_accuracy: 0.9843

Epoch 12/20

469/469 [=====] - 10s 22ms/step - loss: 0.0179 - accuracy: 0.9944 - val_loss: 0.0721 - val_accuracy: 0.9851

Epoch 13/20

469/469 [=====] - 10s 22ms/step - loss: 0.0170 - accuracy: 0.9946 - val_loss: 0.0684 - val_accuracy: 0.9858

```
Epoch 14/20
469/469 [=====] - 11s 23ms/step - loss: 0.0147 -
accuracy: 0.9949 - val_loss: 0.0841 - val_accuracy: 0.9822
Epoch 15/20
469/469 [=====] - 10s 22ms/step - loss: 0.0138 -
accuracy: 0.9956 - val_loss: 0.0802 - val_accuracy: 0.9835
Epoch 16/20
469/469 [=====] - 10s 21ms/step - loss: 0.0139 -
accuracy: 0.9956 - val_loss: 0.0712 - val_accuracy: 0.9849
Epoch 17/20
469/469 [=====] - 10s 22ms/step - loss: 0.0124 -
accuracy: 0.9959 - val_loss: 0.0695 - val_accuracy: 0.9856
Epoch 18/20
469/469 [=====] - 10s 22ms/step - loss: 0.0120 -
accuracy: 0.9960 - val_loss: 0.0772 - val_accuracy: 0.9860
Epoch 19/20
469/469 [=====] - 11s 23ms/step - loss: 0.0094 -
accuracy: 0.9968 - val_loss: 0.0829 - val_accuracy: 0.9851
Epoch 20/20
469/469 [=====] - 9s 20ms/step - loss: 0.0095 -
accuracy: 0.9969 - val_loss: 0.0791 - val_accuracy: 0.9843
```

```
[19]: # Evaluate the model
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
Test loss: 0.07907029986381531
Test accuracy: 0.9843000173568726
```