

lab3-1

April 22, 2025

```
[1]: import tensorflow as tf
      from tensorflow import keras
      from keras import layers
      from keras.datasets import fashion_mnist
```

```
[3]: import ssl
      ssl._create_default_https_context = ssl._create_unverified_context
      # Load the dataset
      (x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-idx1-ubyte.gz>

29515/29515 [=====] - 0s 0us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-idx3-ubyte.gz>

26421880/26421880 [=====] - 7s 0us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx1-ubyte.gz>

5148/5148 [=====] - 0s 0us/step

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx3-ubyte.gz>

4422102/4422102 [=====] - 1s 0us/step

```
[4]: # Normalize the pixel values to be between 0 and 1
      x_train = x_train.astype('float32') / 255
      x_test = x_test.astype('float32') / 255
```

```
[5]: # Convert the labels to one-hot encoded vectors
      num_classes = 10
      y_train = keras.utils.to_categorical(y_train, num_classes)
      y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
[6]: # Build the model
      model = keras.Sequential([
          layers.Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)),
          layers.MaxPooling2D((2,2)),
          layers.Conv2D(64, (3,3), activation='relu'),
          layers.MaxPooling2D((2,2)),
```

```

layers.Flatten(),
layers.Dense(128, activation='relu'),
layers.Dense(num_classes, activation='softmax')
])

```

```

[7]: # Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

```

```

[8]: # Train the model
model.fit(x_train.reshape(-1,28,28,1), y_train, epochs=10, batch_size=32,
         validation_data=(x_test.reshape(-1,28,28,1), y_test))

```

```

Epoch 1/10
1875/1875 [=====] - 42s 22ms/step - loss: 0.4482 -
accuracy: 0.8370 - val_loss: 0.3396 - val_accuracy: 0.8744
Epoch 2/10
1875/1875 [=====] - 40s 21ms/step - loss: 0.3007 -
accuracy: 0.8890 - val_loss: 0.3058 - val_accuracy: 0.8884
Epoch 3/10
1875/1875 [=====] - 40s 21ms/step - loss: 0.2527 -
accuracy: 0.9065 - val_loss: 0.2837 - val_accuracy: 0.8954
Epoch 4/10
1875/1875 [=====] - 40s 21ms/step - loss: 0.2211 -
accuracy: 0.9176 - val_loss: 0.2597 - val_accuracy: 0.9040
Epoch 5/10
1875/1875 [=====] - 40s 21ms/step - loss: 0.1944 -
accuracy: 0.9270 - val_loss: 0.2541 - val_accuracy: 0.9103
Epoch 6/10
1875/1875 [=====] - 40s 21ms/step - loss: 0.1710 -
accuracy: 0.9346 - val_loss: 0.2493 - val_accuracy: 0.9119
Epoch 7/10
1875/1875 [=====] - 41s 22ms/step - loss: 0.1511 -
accuracy: 0.9431 - val_loss: 0.2688 - val_accuracy: 0.9097
Epoch 8/10
1875/1875 [=====] - 42s 22ms/step - loss: 0.1310 -
accuracy: 0.9507 - val_loss: 0.2764 - val_accuracy: 0.9055
Epoch 9/10
1875/1875 [=====] - 41s 22ms/step - loss: 0.1169 -
accuracy: 0.9557 - val_loss: 0.2817 - val_accuracy: 0.9144
Epoch 10/10
1875/1875 [=====] - 41s 22ms/step - loss: 0.1031 -
accuracy: 0.9603 - val_loss: 0.2955 - val_accuracy: 0.9083

```

```

[8]: <keras.callbacks.History at 0x15ae3bb50>

```

```
[9]: # Evaluate the model
test_loss, test_acc = model.evaluate(x_test.reshape(-1,28,28,1), y_test,
↳ verbose=2)
print('Test accuracy:', test_acc)
```

313/313 - 1s - loss: 0.2955 - accuracy: 0.9083 - 1s/epoch - 5ms/step
Test accuracy: 0.90829998254776

0.0.1 THANK YOU